

# sklearn.linear\_model.LinearRegression

```
class sklearn.linear_model.LinearRegression(*, fit_intercept=True, copy_X=True, n_jobs=None, positive=False) \[source\]
```

Ordinary least squares Linear Regression.

LinearRegression fits a linear model with coefficients  $w = (w_1, \dots, w_p)$  to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation.

## Parameters:

**fit\_intercept** : *bool, default=True*

Whether to calculate the intercept for this model. If set to False, no intercept will be used in calculations (i.e. data is expected to be centered).

**copy\_X** : *bool, default=True*

If True, X will be copied; else, it may be overwritten.

**n\_jobs** : *int, default=None*

The number of jobs to use for the computation. This will only provide speedup in case of sufficiently large problems, that is if firstly `n_targets > 1` and secondly X is sparse or if `positive` is set to `True`. `None` means 1 unless in a [joblib.parallel\\_backend](#) context. `-1` means using all processors. See [Glossary](#) for more details.

**positive** : *bool, default=False*

When set to `True`, forces the coefficients to be positive. This option is only supported for dense arrays.

*New in version 0.24.*

## Attributes:

**coef\_** : *array of shape (n\_features, ) or (n\_targets, n\_features)*

Estimated coefficients for the linear regression problem. If multiple targets are passed during the fit (y 2D), this is a 2D array of shape (n\_targets, n\_features), while if only one target is passed, this is a 1D array of length n\_features.

**rank\_** : *int*

Rank of matrix X. Only available when X is dense.

**singular\_** : *array of shape (min(X, y),)*

Singular values of X. Only available when X is dense.

**intercept\_** : *float or array of shape (n\_targets,)*

Independent term in the linear model. Set to 0.0 if `fit_intercept = False`.

**n\_features\_in\_** : *int*

Number of features seen during [fit](#).

*New in version 0.24.*

**feature\_names\_in\_** : *ndarray of shape (n\_features\_in\_,)*

Names of features seen during [fit](#). Defined only when X has feature names that are all strings.

*New in version 1.0.*

## See also:

[Ridge](#)

Ridge regression addresses some of the problems of Ordinary Least Squares by imposing a penalty on the size of the coefficients with l2 regularization.

### Lasso

The Lasso is a linear model that estimates sparse coefficients with l1 regularization.

### ElasticNet

Elastic-Net is a linear regression model trained with both l1 and l2 -norm regularization of the coefficients.

## Notes

From the implementation point of view, this is just plain Ordinary Least Squares (`scipy.linalg.lstsq`) or Non Negative Least Squares (`scipy.optimize.nnls`) wrapped as a predictor object.

## Examples

```
>>> import numpy as np
>>> from sklearn.linear_model import LinearRegression
>>> X = np.array([[1, 1], [1, 2], [2, 2], [2, 3]])
>>> # y = 1 * x_0 + 2 * x_1 + 3
>>> y = np.dot(X, np.array([1, 2])) + 3
>>> reg = LinearRegression().fit(X, y)
>>> reg.score(X, y)
1.0
>>> reg.coef_
array([1., 2.])
>>> reg.intercept_
3.0...
>>> reg.predict(np.array([[3, 5]]))
array([16.] )
```

## Methods

<a href="#"><code>fit(X, y[, sample_weight])</code></a>	Fit linear model.
<a href="#"><code>get_metadata_routing()</code></a>	Get metadata routing of this object.
<a href="#"><code>get_params([deep])</code></a>	Get parameters for this estimator.
<a href="#"><code>predict(X)</code></a>	Predict using the linear model.
<a href="#"><code>score(X, y[, sample_weight])</code></a>	Return the coefficient of determination of the prediction.
<a href="#"><code>set_fit_request(*[, sample_weight])</code></a>	Request metadata passed to the <code>fit</code> method.
<a href="#"><code>set_params(**params)</code></a>	Set the parameters of this estimator.
<a href="#"><code>set_score_request(*[, sample_weight])</code></a>	Request metadata passed to the <code>score</code> method.

**`fit(X, y, sample_weight=None)`**

[\[source\]](#)

Fit linear model.

**Parameters:**

**X : {array-like, sparse matrix} of shape (n\_samples, n\_features)**

Training data.

**y : array-like of shape (n\_samples,) or (n\_samples, n\_targets)**

Target values. Will be cast to X's dtype if necessary.

**sample\_weight : array-like of shape (n\_samples,), default=None**

Individual weights for each sample.

*New in version 0.17:* parameter *sample\_weight* support to LinearRegression.

**Returns:**

**self : object**

Fitted Estimator.

**get\_metadata\_routing()**[\[source\]](#)

Get metadata routing of this object.

Please check [User Guide](#) on how the routing mechanism works.

**Returns:**

**routing : MetadataRequest**

A [MetadataRequest](#) encapsulating routing information.

**get\_params(deep=True)**[\[source\]](#)

Get parameters for this estimator.

**Parameters:**

**deep : bool, default=True**

If True, will return the parameters for this estimator and contained subobjects that are estimators.

**Returns:**

**params : dict**

Parameter names mapped to their values.

**predict(X)**[\[source\]](#)

Predict using the linear model.

**Parameters:**

**X : array-like or sparse matrix, shape (n\_samples, n\_features)**

Samples.

**Returns:**

**C : array, shape (n\_samples,)**

Returns predicted values.

**score(X, y, sample\_weight=None)**[\[source\]](#)

Return the coefficient of determination of the prediction.

The coefficient of determination  $R^2$  is defined as  $(1 - \frac{u}{v})$ , where  $u$  is the residual sum of squares  $((y_{\text{true}} - y_{\text{pred}})**2).sum()$  and  $v$  is the total sum of squares  $((y_{\text{true}} - y_{\text{true}.mean()})**2).sum()$ . The best possible score is 1.0 and it can be negative (because the model can be arbitrarily worse). A constant model that always predicts the expected value of  $y$ , disregarding the input features, would get a  $R^2$  score of 0.0.

#### Parameters:

**X : array-like of shape (n\_samples, n\_features)**

Test samples. For some estimators this may be a precomputed kernel matrix or a list of generic objects instead with shape  $(n_{\text{samples}}, n_{\text{samples\_fitted}})$ , where  $n_{\text{samples\_fitted}}$  is the number of samples used in the fitting for the estimator.

**y : array-like of shape (n\_samples,) or (n\_samples, n\_outputs)**

True values for X.

**sample\_weight : array-like of shape (n\_samples,), default=None**

Sample weights.

#### Returns:

**score : float**

$R^2$  of `self.predict(X)` w.r.t. `y`.

#### Notes

The  $R^2$  score used when calling `score` on a regressor uses `multioutput='uniform_average'` from version 0.23 to keep consistent with default value of `r2_score`. This influences the `score` method of all the multioutput regressors (except for [MultiOutputRegressor](#)).

**set\_fit\_request**(\*, sample\_weight: [Union\[bool, None, str\]](#) = '\$UNCHANGED\$') → [LinearRegression](#)

[\[source\]](#)

Request metadata passed to the `fit` method.

Note that this method is only relevant if `enable_metadata_routing=True` (see [sklearn.set\\_config](#)). Please see [User Guide](#) on how the routing mechanism works.

The options for each parameter are:

- `True`: metadata is requested, and passed to `fit` if provided. The request is ignored if metadata is not provided.
- `False`: metadata is not requested and the meta-estimator will not pass it to `fit`.
- `None`: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
- `str`: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (`sklearn.utils.metadata_routing.UNCHANGED`) retains the existing request. This allows you to change the request for some parameters and not others.

*New in version 1.3.*

**Note:** This method is only relevant if this estimator is used as a sub-estimator of a meta-estimator, e.g. used inside a [Pipeline](#). Otherwise it has no effect.

#### Parameters:

**sample\_weight : str, True, False, or None, default=sklearn.utils.metadata\_routing.UNCHANGED**

Metadata routing for `sample_weight` parameter in `fit`.

#### Returns:

**self : object**

The updated object.

**set\_params(\*\*params)**[\[source\]](#)

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as [Pipeline](#)). The latter have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

**Parameters:****\*\*params** : *dict*

Estimator parameters.

**Returns:****self** : *estimator instance*

Estimator instance.

**set\_score\_request(\*, sample\_weight: [Union\[bool, None, str\]](#) = '\$UNCHANGED\$') → [LinearRegression](#)**[\[source\]](#)

Request metadata passed to the `score` method.

Note that this method is only relevant if `enable_metadata_routing=True` (see [sklearn.set\\_config](#)). Please see [User Guide](#) on how the routing mechanism works.

The options for each parameter are:

- `True`: metadata is requested, and passed to `score` if provided. The request is ignored if metadata is not provided.
- `False`: metadata is not requested and the meta-estimator will not pass it to `score`.
- `None`: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
- `str`: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (`sklearn.utils.metadata_routing.UNCHANGED`) retains the existing request. This allows you to change the request for some parameters and not others.

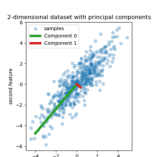
*New in version 1.3.*

**Note:** This method is only relevant if this estimator is used as a sub-estimator of a meta-estimator, e.g. used inside a [Pipeline](#). Otherwise it has no effect.

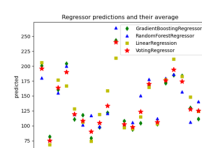
**Parameters:****sample\_weight** : *str, True, False, or None, default=sklearn.utils.metadata\_routing.UNCHANGED*Metadata routing for `sample_weight` parameter in `score`.**Returns:****self** : *object*

The updated object.

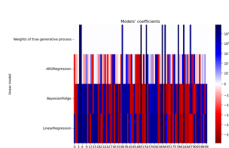
## Examples using `sklearn.linear_model.LinearRegression`



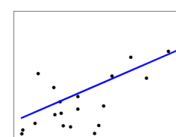
Principal Component



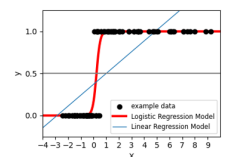
Plot individual and



Comparing Linear

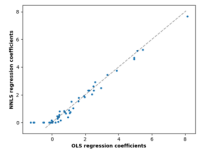


Linear Regression



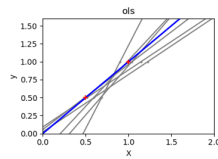
Logistic function

## Regression vs Partial Least Squares Regression



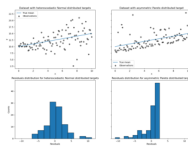
Non-negative least squares

## voting regression predictions



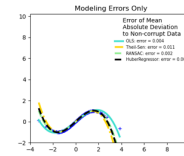
Ordinary Least Squares and Ridge Regression Variance

## Bayesian Regressors

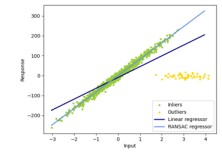


Quantile regression

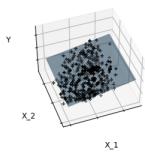
## Example



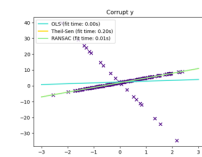
Robust linear estimator fitting



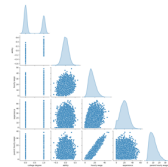
Robust linear model estimation using RANSAC



Sparsity Example: Fitting only features 1 and 2



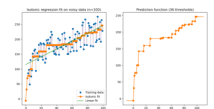
Theil-Sen Regression



Failure of Machine Learning to infer causal effects



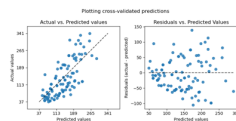
Face completion with a multi-output estimators



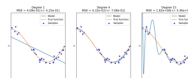
Isotonic Regression



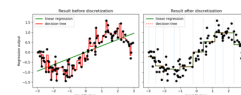
Metadata Routing



Plotting Cross-Validated Predictions



Underfitting vs. Overfitting



Using KBinsDiscretizer to discretize continuous features

Toggle Menu

© 2007 - 2023, scikit-learn developers (BSD License). [Show this page source](#)